

DO: pick a cheap model first.

Your default model matters more than you think. if you didn't change it during setup, check what you're running:

bash

```
openclaw config get agents.defaults.model
```

If it says Opus anywhere, switch immediately. opus is \$5/\$25 per million tokens. sonnet does 90% of the same work at \$3/\$15. For your first week of learning, even cheaper models work fine. GLM-5.1 at \$0.95/\$3.15 or openrouter free tier costs literally nothing.

Someone I helped was spending \$47/week without realizing it. changed one setting. Next week cost \$6.

DON'T: skip the gateway security.

If you're on a VPS or any internet-connected machine:

bash

```
openclaw config get | grep -E "host|bind"
```

If it says 0.0.0.0 Your agent is accessible to anyone who finds your IP. there are 500,000+ OpenCLAW instances exposed on the public internet right now. one had a zero-click exploit (CVE-2026-25253, patched) that let attackers hijack agents from a single webpage visit.

bash

```
openclaw config set gateway.bind loopback
```

two minutes. do it before connecting any channel.

DO: write a SOUL.md with boundaries, not just personality.

Most guides tell you to write personality rules. "be direct, match my tone, don't say absolutely." that's fine. but the part people skip is boundaries:

markdown

```
Never send emails, messages, or make bookings without showing me first.  
Never sign up for services without my explicit approval.  
Never delete files or emails without asking.
```

Without boundaries, your agent will do exactly what it thinks you want at machine speed with zero hesitation. someone told their agent to "explore what you can do." it created dating profiles using data from his emails. the agent wasn't broken. the instructions were too open.

"Never do X" works better than "try to be Y." your SOUL.md is built through irritation, not planning.

DON'T: install skills in your first week.

I know. Clawhub has 13,000+ skills and they all look cool. don't.

1,467 malicious skills were found during the ClawHavoc campaign. typosquatted names, clean documentation, legitimate-looking publishers. they silently exfiltrated your .env file (every API key, every OAuth token) to external servers.

clawhub has VirusTotal scanning now. It's better. But "better" and "safe" are different things.

learn what your agent can do without skills first. You'll be surprised. After week 1, add ONE skill from a verified publisher. test it for a few days. Watch your costs. then add another. never more than one at a time.

DO: use /new aggressively.

Every message you send in a session gets included in every future API call. after a few days of chatting, you're sending thousands of tokens of old conversation with every new message. that costs money and makes your agent slower and more confused.

`/new` starts a fresh session. Your agent keeps all its memory files, SOUL.md, everything. You're just clearing the conversation buffer.

Use it before any big task. When your agent starts acting weird. at least once a day as a habit.

also learn `/btw` for tangent questions. Instead of polluting your main session with "what's the weather tomorrow," type `/btw what's the weather tomorrow` and it fires off a side conversation without touching your main context.

DON'T: create a second agent.

Every new user thinks they need multiple agents. personal, work, coding. you don't. not yet.

Every agent is an independent token consumer. every agent needs its own channel binding. Every agent complicates debugging. i've seen too many people create a second agent to "fix" problems with the first one. now they have two broken agents.

Get one agent working perfectly for 2 weeks. then decide if you actually need another. most people don't.

DO: check your costs every single day for the first 2 weeks.

check your API provider's dashboard directly (console.anthropic.com, platform.openai.com, whatever you use). Don't rely on OpenClaw's internal cost tracking, it's an estimate and sometimes doesn't match what you actually get billed.

on Sonnet with one agent and no skills, expect \$3-8/month for moderate personal use. if you're above that in your first week, something is wrong and it's fixable.

Watch for heartbeat costs specifically. OpenClaw checks in every 30-60 minutes. if those heartbeats are running on your expensive model, you're paying for your agent to check its own pulse 24 times a day at premium rates.

DON'T: auto-update without checking the changelog.

This is the mistake experienced users make. OpenClaw updates 2-3 times a week. Some updates break things. if you auto-update overnight, you might wake up to a broken setup with no idea what changed.

Either pin your version and update manually when you're ready, or at minimum read the changelog before letting updates through.

DO: have realistic expectations for your first week.

day 1-2: set up your model, lock your gateway, write your SOUL.md. have normal conversations. ask stupid questions. get comfortable.

day 3-4: start using it for real tasks. calendar, reminders, web searches, summarizing articles. the boring stuff. keep everything read-only. Don't give it write access to email or files yet.

day 5-7: refine your SOUL.md based on what annoyed you. Check your costs. get a feel for daily usage.

That's it. no skills. no second agent. no multi-agent orchestrator. no cron jobs. just one agent that knows who you are, respects boundaries, and does basic tasks reliably.

If that feels underwhelming, good. The people still using OpenCLaw three months from now all started exactly like this. The people who quit started with 8 agents and 20 skills on day one. For more details, you can visit [r/better_claw](https://www.reddit.com/r/better_claw).